

New upper bounds on covering codes $K_q(n, R)$ for alphabets of size six and seven

Mark Marosi* and Claude (Anthropic)[†]

August 2026

Abstract

We present improved upper bounds for nine entries of the standard tables of bounds on $K_q(n, R)$, the minimum cardinality of a q -ary code of length n with covering radius R , for $q \in \{6, 7\}$: $K_6(7, 3) \leq 233$, $K_6(8, 3) \leq 1045$, $K_6(8, 4) \leq 167$, $K_6(9, 4) \leq 710$, $K_6(9, 5) \leq 125$, $K_6(10, 4) \leq 2951$, $K_6(10, 5) \leq 610$, $K_7(8, 4) \leq 331$, and $K_7(9, 4) \leq 1814$. The previous best bounds, recorded in Kéri’s tables (last updated 2011), all arose from general constructions (direct sums and related product rules) rather than from explicit search; to our knowledge these are the first improvements to any upper bound on $K_q(n, R)$ with $q \geq 5$ since 2011. The new bounds were found by focused local search seeded with the construction-based incumbents. All nine codes are given explicitly in the ancillary files, together with a standalone verifier; each code was checked by four independent exhaustive verification methods.

1 Introduction

Let $\mathbb{Z}_q^n$ denote the Hamming space of q -ary words of length n . A code $C \subseteq \mathbb{Z}_q^n$ has *covering radius* R if every word of $\mathbb{Z}_q^n$ is within Hamming distance R of some codeword. The function

$$K_q(n, R) = \min\{|C| : C \subseteq \mathbb{Z}_q^n \text{ has covering radius } \leq R\}$$

is a central quantity of covering-code theory [1]; the binary case $K_2(n, 1)$ includes the domination number of the hypercube, and the ternary case $K_3(n, 1)$ is the classical football pool problem.

Upper bounds on $K_q(n, R)$ are established by explicit constructions, and the state of the art is recorded in the tables of Kéri [2], which subsume and extend the earlier tables of Cohen, Honkala, Litsyn and Lobstein [1]. The last update to Kéri’s tables was made in November 2011. Recent activity on $K_q(n, R)$ has concentrated on *lower* bounds: Wu and Chen [4] improved binary lower bounds via congruence arguments, and Gijswijt and Polak [3] obtained many new lower bounds for $q \leq 5$ from semidefinite programming. On the upper-bound side, for alphabets $q \geq 5$ we are not aware of any published improvement since the 2011 freeze; a systematic search of the literature (arXiv, DBLP, OpenAlex, Lobstein’s covering-code bibliography, and the publication lists of the authors active in this area) found none.

*Department of Measurement and Information Systems, Budapest University of Technology and Economics (BME), Budapest, Hungary. email: marosi@mit.bme.hu.

[†]The search software, the verification protocol, the constructions, and this manuscript were produced autonomously by the AI system Claude (Anthropic) on a single NVIDIA GH200 node; M.M.’s contribution was the idea of applying large-scale GPU/CPU search to the frozen covering-code tables and providing the computational resources. See Section 3.

The reason the entries treated here were improvable is visible in Kéri’s own source annotations. For moderate q and $R \geq 3$ nearly all tabulated upper bounds carry keys denoting *general constructions*: the direct sum

$$K_q(n_1 + n_2, R_1 + R_2) \leq K_q(n_1, R_1) \cdot K_q(n_2, R_2), \quad (1)$$

the special case $K_q(n + 1, R) \leq q K_q(n, R)$, the monotonicity $K_q(n + 1, R + 1) \leq K_q(n, R)$, and product rules over alphabet factorizations. These are inherited arithmetic: no search was ever run in those cells. The present note shows that plain local search, seeded with the construction that produced the incumbent, beats nine such bounds in the range $q \in \{6, 7\}$, $n \leq 10$ — in one case by 23%.

2 Results

cell	previous bounds	new UB	Δ	code file	previous UB from
$K_6(7, 3)$	70–246	233	−13	K6_7_3_M233.txt	$6 \cdot K_6(6, 3)$
$K_6(8, 3)$	246–1080	1045	−35	K6_8_3_M1045.txt	$K_6(4, 1) \cdot K_6(4, 2)$
$K_6(8, 4)$	46–216	167	−49	K6_8_4_M167.txt	$6 \cdot K_6(7, 4)$
$K_6(9, 4)$	136–738	710	−28	K6_9_4_M710.txt	$K_6(3, 1) \cdot K_6(6, 3)$
$K_6(9, 5)$	32–144	125	−19	K6_9_5_M125.txt	alphabet product rule
$K_6(10, 4)$	417–2952	2951	−1	K6_10_4_M2951.txt	$K_6(4, 1) \cdot K_6(6, 3)$
$K_6(10, 5)$	83–615	610	−5	K6_10_5_M610.txt	$K_6(6, 3) \cdot K_6(4, 2)$
$K_7(8, 4)$	76–343	331	−12	K7_8_4_M331.txt	$K_7(7, 3)$ (monotonicity)
$K_7(9, 4)$	264–1843	1814	−29	K7_9_4_M1814.txt	direct sum

Table 1: New upper bounds. “Previous bounds” are the lower–upper bound pairs of Kéri’s tables [2]; lower bounds are unaffected by this work. The last column gives the construction from which the previous upper bound was inherited, per the source keys of [2].

The nine codes are supplied as ancillary files (one codeword per line, digits over $\{0, \dots, q-1\}$). None of the searches had converged when stopped, so we expect several of these bounds to be improvable further by the same method with more computation.

Remark 1. *An audit of the tables of [2] against their own derivation rules found one internal gap: the tabulated $K_{17}(7, 2) \leq 252735$ is weaker than $K_{17}(7, 2) \leq 17 \cdot K_{17}(6, 2) = 17 \cdot 14424 = 245208$, which follows from the tabulated $K_{17}(6, 2) \leq 14424$ and the standard rule $K_q(n+1, R) \leq q K_q(n, R)$. We record this as an observation; it involves no new construction.*

3 Method

3.1 Search

We search for a code of prescribed size M minimizing the number of uncovered words. The state is a multiset of M codewords together with the coverage count $\text{cnt}(w)$ of every $w \in \mathbb{Z}_q^n$; a move changes one coordinate of one codeword. The move update is cheap because of the following observation.

Lemma 1. *Let c' be obtained from c by setting coordinate p to $v \neq c_p$. Then the set of words leaving the radius- R ball of c and the set of words entering the ball of c' are both indexed by the words at distance exactly R from c in the coordinates other than p ; consequently a move updates exactly $2S$ coverage counters with $S = \binom{n-1}{R}(q-1)^R$, the two sets are disjoint, and distinct counter updates address distinct words.*

Proof. Since $d(w, c') = d(w, c) + [w_p = c_p] - [w_p = v]$, a word can leave the ball of c only if $d(w, c) = R$ and $w_p = c_p$, and can enter the ball of c' only if $d(w, c') = R$ and $w_p = v$; writing such a word as its restriction to the coordinates $\neq p$ (at distance exactly R from the restriction of c) plus its value at p gives the stated bijections, disjointness ($w_p = c_p$ versus $w_p = v$), and injectivity. $\square$

The search itself is a focused local search in the WalkSAT lineage, close in spirit to the tabu searches that produced much of the original tables [6, 5]: pick a random uncovered word u ; collect as candidate moves every codeword at distance exactly $R + 1$ from u nudged one coordinate towards u (if no codeword is that close, the codewords at minimum distance from u are used); evaluate all candidates exactly and in parallel; commit a best candidate, ties broken uniformly at random. On stagnation a codeword is teleported onto an uncovered word. Parameter studies on $K_6(6, 3)$ showed that *candidate breadth dominates*: evaluating all candidate moves exactly outperformed sampling 24 or 64 of them (final uncovered counts 0 versus 82 and 18), while tabu tenures and WalkSAT-style noise were harmful on these instances. Simulated annealing sustained a $33\times$ higher move rate and still lost to the focused search — move quality beats move count here.

3.2 Seeding and descent

For each cell we first rebuilt a code of the incumbent size from the same construction that produced the incumbent bound (e.g., for $K_6(10, 4)$ the direct sum (1) of an optimal $K_6(4, 1)$ code of size 72 and a $K_6(6, 3)$ code of size 41, both found by direct search and verified). Descent then proceeds by remove-and-repair: delete the codeword whose private coverage (words covered by no other codeword) is smallest, and repair the deficit by local search; on success, repeat at $M - 1$. Searching the small factors and re-composing (a “factor attack”) was decisive for seeding; the improvements themselves come from the descent breaking out of the product structure.

3.3 Hardware and effort

All record searches ran on the 64-core Grace CPU of a single NVIDIA GH200 node (a CUDA port was implemented and verified but was latency-bound and slower than the CPU at these instance sizes). Per-cell effort ranged from seconds ($K_6(8, 4)$, $K_6(9, 5)$) to under two hours of wall clock; none of the descents had converged when stopped. Move throughput ranged from $\sim 1.3\text{k}$ moves/s per worker on the largest cells to $\sim 400\text{k}$ on the smallest.

4 Verification

False records are worse than no records, so verification is exhaustive, exact, and redundant. Every claimed code was checked by four methods, two of which share no code with the other two:

1. ball marking: mark all words within distance R of each codeword in a q^n bitmap and confirm the bitmap is full;
2. meet-in-the-middle bitset covering over a coordinate split;
3. a min-distance scan: enumerate all q^n words and compute each word’s distance to the nearest codeword (this also reports the exact covering radius);
4. Hamming-graph dilation: R -fold neighborhood dilation of the code’s indicator vector, computing no distances at all, with cost independent of M .

In addition, a fifth, independently written dilation verifier (produced without access to the other implementations) re-confirmed all nine codes from the published files. The largest check ($q^n = 6^{10} = 60,466,176$ words) completes in seconds.

The ancillary files include `verify_cov.py`, a standalone dependency-free Python verifier: `python3 verify_cov.py CODEFILE q n R` exhaustively re-checks any of the claimed bounds in at most a few minutes.

5 Discussion

The pattern in our experiments is clean: local search consistently beat bounds that were inherited from general constructions (relative slack between the previous upper bound and the sphere-covering lower bound of $4\times\text{--}10\times$), and consistently failed to improve cells that had already been the object of dedicated search or algebraic construction (slack $\leq 3\times$), such as $K_3(11, 4) \leq 81$, $K_6(6, 3) \leq 41$, and $K_8(8, 4) \leq 512$. Kéri’s tables contain many further cells of the first kind, for $q \in \{6, \dots, 21\}$ and $R \geq 3$; we expect them to fall to the same method, and a systematic sweep is in progress. Since the tables have had no maintainer since 2011, we also intend to publish a merged, machine-readable bounds table (Kéri 2011 plus the post-2011 lower-bound literature plus the present improvements) alongside the code.

Data availability

The codes, the verifier, and a README are ancillary files of this arXiv submission. The full search and verification source code is public at <https://github.com/Mapika/coldcase> (covering-code track under `cov/`).

Acknowledgments

This work is AI-assisted in the strong sense: the discovery and verification pipeline was designed, implemented, and operated autonomously by Claude (Anthropic). The human author’s contributions were the initiating idea (applying modern large-scale search hardware to record tables frozen in the 2000s) and the computational resources.

References

- [1] G. Cohen, I. Honkala, S. Litsyn, A. Lobstein, *Covering Codes*, North-Holland, Amsterdam, 1997.

- [2] G. Kéri, Tables for bounds on covering codes, <https://old.sztaki.hu/~keri/codes/> (last updated November 2011).
- [3] D. Gijswijt, S. Polak, Semidefinite lower bounds for covering codes, arXiv:2504.01932, 2025.
- [4] Y. Wu, C. Chen, Improved lower bounds on the domination number of hypercubes and binary codes with covering radius one, Discrete Mathematics 347 (2024) 113752.
- [5] P. J. M. van Laarhoven, E. H. L. Aarts, J. H. van Lint, L. T. Wille, New upper bounds for the football pool problem for 6, 7 and 8 matches, J. Combin. Theory Ser. A 52 (1989) 304–312.
- [6] P. R. J. Östergård, Constructing covering codes by tabu search, J. Combin. Des. 5 (1997) 71–80.
- [7] G. Kéri, P. R. J. Östergård, Further results on the covering radius of small codes, Discrete Mathematics 307 (2007) 69–77.
- [8] A. Florath, A Lean 4 formalization of q -ary covering codes with a proof-carrying database of certified bounds, arXiv:2606.09600, 2026.

A The code proving $K_6(8, 4) \leq 167$

For self-containedness we print the largest single improvement ($216 \rightarrow 167$, -23%). Each row lists codewords of $\mathbb{Z}_6^8$; the machine-readable version is the ancillary file `K6_8_4_M167.txt`.

```

42030014 50143021 30440244 25403541
40250502 02441030 35400150 50023554
50225151 12333122 33515545 44121521
14212440 42314041 35042002 34355041
05025020 12253354 05133334 42335434
03201411 35343410 24143214 54542204
32151400 10440535 43332055 32200053
30051122 00045454 52024303 50551345
15230201 24410415 54434031 23142140
33254504 25214034 21244445 53231213
04130543 01302114 34325315 50135422
45510313 51220320 45504535 22125553
31533321 55255112 04041551 42150315
02322452 15142345 02303401 04013213
33310230 11101302 41241312 25550524
13423003 22044222 20311131 14551115
02105344 10310252 10344333 42540351
22520011 15221421 51124410 24221205
11152242 01253131 04202301 53440342
11031550 20131443 30014311 22432502
05111025 41054323 33130421 00504200
20532144 41122233 34544402 40202220
21401320 23553251 51540443 55412301
53051015 32234252 12455213 12000145
31105501 01212555 40424455 13501210
52103235 35321243 24030300 24100153
05431103 33343524 31114105 02214523
45311104 40401434 53352125 01424040
21352033 30102012 03250042 52511554
43120104 43112350 20205330 34452453
54355500 04545025 23535224 44405124
11015404 50015125 50313003 01310221

```

41321232	00553430	25032211	51342511
14115532	52513140	55154250	43054444
23323342	31441054	33022530	02422344
12502423	34233035	00434510	41505000
30521513	05524132	04301325	21413452
15533052	44344100	14052131	40033105
25355405	11335355	41035241	13444111
14324024	23004353	03455235	33245133
44453250	54304042	43512332	13003512
31413225	55000434	45243543	